


Introducing Gridwright

It devours a Figma page and
returns it built: measured
against the design, every
piece filed where your
codebase keeps it.

98.54%

Fidelity at 1440px, a whole landing page.
12 sections, one command.



A prompt is a suggestion

The obvious move is a long prompt that tells an agent how to build a layout. I tried it. An earlier repo had a five-phase workflow written in prose, and the agent skipped the analysis every time the request looked simple.

So the workflow stopped being text.

Claude doesn't pick the next step. It asks.


```
gw next --json → do exactly that stage
→ gw done → gw next --json → ...
```

The workflow is a state machine on disk, and a CLI enforces it.


CODE DOES THIS	THE MODEL DOES THIS
Pull the node, the assets and the reference	Write the component the way this repo writes them
Distill the tree into an IR	Name things, decide the prop API
Match and classify tokens	Name the new tokens
Index existing components	Decide what to reuse
Render, measure, diff	Read the diff and fix it

If it can be checked with an assert, the model doesn't do it.

The Figma, then what it built



FIGMA



GRIDWRIGHT • 1440PX • 91.73

Same band. The headline is in the wrong typeface: see 07.

Twelve sections. Twelve agents.

gw build --view takes each child of the page as a section, runs one sub-agent per section in parallel, and composes the page.

SECTION	SCORE	SECTION	SCORE
NavbarDefault layout	95.70	TabsDefault	95.32
HeroDefault	91.73	Testimonials	97.28
Logos	85.81	PricingDefault	97.78
BentoGrid2x2 *	95.43	FAQDefault	100.00
ItemsDefault	98.23	CTA	98.81
FeatureRising	88.87	Footer layout	97.23

Composed page: 98.54% · structural 99.47 · chromatic 100 · perceptual 95.22

* Rejected by distill first, built on its own later. See 07.

80 values in the design. 65 were already there.

Resolved before a line of the component is written. The new ones are asked once, for the whole page.

65

exact — already in the system, used as is

7

near — within ΔE 1 or a couple of pixels, the system's value wins


8

new — needs a name and a decision

Against 38 project tokens + 513 from Tailwind's theme

Otherwise the model writes `bg-[#1a1a1a]`
and someone refactors.

99.8% against 89%



One section of Untitled UI, built twice from the same commit, in two clones that never saw each other. Same prompt, same ruler. One arm had Gridwright.

The control laid the mosaic out with **position: absolute**.

It got the layout right, 94% structural. It lost on placeholder photos and an approximate brand colour, a minute's fix each. `position: absolute` isn't: it breaks at the first width the design wasn't drawn for.

The score didn't see the typeface

The builds came out in the project's default font instead of Inter, and the score didn't notice: the perceptual diff masks text. Now `gw build` warns when the project doesn't load a typeface from the design. It doesn't substitute and it doesn't install.

Distill rejected one Launch UI section: layers in absolute position, no auto-layout, nothing to extract. The page went on with the other eleven, and that rejection surfaced a bug, the view's registration stopped at the skipped section. Fixed, then the section was built on its own.

This frame does not use auto-layout, so there is no layout to infer. A better prompt will not fix this: it gets fixed in Figma.

Responsive is next

Every score here is at the width the design was drawn at. At 375 and 768 the same page gives 63% and 65%: those widths have no frame to measure against. No fix yet.

STARTED	3 September 2026
COMMITTS	64
TESTS	250
RUNTIME DEPS	5
SPEC	10 laws, written first
INSTALL	Repo only, no npm yet

Open source, MIT.
Link in the first comment.

